

Automatically Synthesizing SQL Queries from Input-Output Examples

Sai Zhang Yuyin Sun
Computer Science & Engineering
University of Washington, USA
{szhang, sunyuyin}@cs.washington.edu

Abstract—Many computer end-users, such as research scientists and business analysts, need to frequently query a database, yet lack enough programming knowledge to write a correct SQL query. To alleviate this problem, we present a *programming by example* technique (and its tool implementation, called SQLSynthesizer) to help end-users automate such query tasks. SQLSynthesizer takes from users an example input and output of how the database should be queried, and then synthesizes a SQL query that reproduces the example output from the example input. Later, when the synthesized SQL query is applied to another, potentially larger, database with a similar schema as the example input, the synthesized SQL query produces a corresponding result that is similar to the example output.

We evaluated SQLSynthesizer on XXX SQL exercises from a classic database textbook and XXX SQL questions raised by real-world users from online forums. SQLSynthesizer synthesizes correct queries for XXX textbook exercises and XXX forum questions, and it does so with small input-output examples.

I. INTRODUCTION

The big data revolution over the past few years has resulted in significant advances in digitization of massive amounts of data and accessibility of computational devices to massive proportions of the population. A key challenge faced by many enterprise or computer end-users nowadays is the management of their increasingly large and complex databases.

Motivation. Although the relational database management system (RDBMS) and the de facto language (SQL) are perfectly adequate for most end-users’ needs [15], the costs associated with deployment and use of database software and SQL are prohibitive. For example, as pointed out in [9], conventional RDBMS software remains underused in the long tail of science: the large number of users, such as the research scientists who are in relatively small labs and individual researchers, have limited IT training, staff and infrastructure yet collectively produce the bulk of scientific knowledge.

The problem is exacerbated by the fact that many end-users have myriad diverse backgrounds including business analysts, commodity traders, human resource managers, finance professionals, and marketing managers. Those end-users are not professional programmers, but are experts in some other domains. They need to retrieve a variety of information from their database and use the information to support their business decisions. Although most end-users can clearly describe *what* the task is, they are often stuck with the process of *how* to write a correct database query (i.e., a SQL query) even after receiving

step-by-step, detailed, and syntactically correct instructions. Thus, typical end-users often need to seek information from online help forums, or ask SQL experts. This process can be repetitive, laborious, and frustrating. To assist non-expert end-users in conducting database query tasks, a highly accessible tool that can be used to “describe” their needs and “connect” their intentions to executable SQL queries would be highly desirable.

Existing Solutions. *Graphical User Interfaces* (GUIs) and *general programming languages* are two state-of-the-art approaches in helping end-users perform database queries. However, both approaches are far from satisfactory.

Many RDBMS come with a well-designed GUI with tons of features. However, a GUI is often fixed, and does not permit users to personalize a database’s functionality for their query tasks. On the other hand, as a GUI supports more and more customization features, users may struggle to discover those features, which can significantly degrade its usability.

General programming languages, such as SQL, Java (with JDBC), or other domain specific query languages, serve as a fully expressive medium for communicating a user’s intention to a database. However, general purpose programming languages have never been easy for end-users who are not professional programmers. Learning a practical programming language (even a simplified, high-level domain specific language, such as MDX [20]) often requires a substantial amount of time and energy that a typical end-user would not prefer, and should not be expected, to invest.

Our Solution: Synthesizing SQL Queries from Input-Output Examples In this paper, we present a technique (and its tool implementation, called SQLSynthesizer) to automatically synthesize SQL queries from input-output examples¹. Although input-output examples may lead to underspecification, writing them, as opposed to writing declarative specifications or imperative code of any form, is one of the most straightforward ways to describe *what* the task is and is more natural for a typical end-user to use. If the synthesized SQL query is applied to the example input, then it produces the example output; and if the SQL query is applied to other similar inputs (potentially much larger tables), then the SQL query produces a corresponding output.

¹The SQL queries in this paper refer to the read-only database queries, which do not modify the database content.

SQLSynthesizer is designed to be used by non-expert database end-users when they do not know how to write a correct SQL query. End-users can use SQLSynthesizer to obtain a SQL query to transform multiple, huge database tables by constructing small, representative input and output example tables. We also envision SQLSynthesizer to be useful in an online education setting (i.e., an online database course). Recently, several education initiatives such as EdX, Coursera, and Udacity are teaming up with experts to provide high quality online courses to several thousands of students worldwide. One challenge, which is not present in a traditional classroom setting, is to provide answers on questions raised by a large number of students. A tool, like SQLSynthesizer, that has the potential of answer SQL query related questions would be useful.

Inferring SQL queries from examples is challenging, primarily for two reasons. First, the standard SQL language is inherently complex; a SQL query can consist of many parts, such as joins, aggregates, the `GROUP BY` clause, and the `ORDER BY` clause. Searching for a SQL query to satisfy the given input-output examples, as proved by Sarma et al. [4], is a PSPACE-hard problem. Thus, a brute-force approach such as exhaustively enumerating all syntactically-valid SQL queries and then filtering away those do not satisfy the examples would quickly become intractable in practice. Second, a SQL query has a rich set of operations: it needs to be evaluated on *multiple* input tables; it needs to perform data grouping, selection, and ordering; and it needs to project data on certain columns as the output. All such operations must be handled properly.

To make SQL query synthesis from examples feasible in practice, our SQLSynthesizer technique focuses on a widely-used SQL subset (Section III), and uses three steps to link a user's intention to a desirable SQL query:

- **Skeleton Creation.** SQLSynthesizer scans the given input-output examples and heuristically identify the table joins and projection columns in the result query. Then, it creates an incomplete SQL query (called, query skeleton) to capture the basic structure of the result query.
- **Query Completion.** SQLSynthesizer uses a rule-learning algorithm from the machine learning community to infer a set of accurate and expressive rules, which transform the input example into the output example. Then, it searches for possible aggregates and columns in the `ORDER BY` clause; and outputs a list of syntactically-valid candidate queries.
- **Candidate Ranking.** If multiple SQL queries satisfy the given input-output examples, SQLSynthesizer employs the Occam's razor principle to rank more likely queries higher in the output.

Compared to previous approaches [3], [4], [26], [28], SQLSynthesizer has two notable features:

- **It is fully automated.** Besides an example input and output pair, SQLSynthesizer does not require users to provide annotations or hints of any form. This distinguishes our work from competing techniques such as specification-based query inference [28] and query synthesis from

imperative code [3].

- **It supports a wide range of SQL queries.** Similar approaches in the literature support a small subset of the SQL language; most of them can only infer simple select-from-where queries on a single table [3], [4], [4], [26], [28]. By contrast, SQLSynthesizer significantly enriches the supported SQL subset. Besides supporting the standard select-from-where queries, SQLSynthesizer also supports many other important SQL features, such as table joins, aggregates (e.g., `MAX()`, `MIN()`, `SUM()`, and `COUNT()`), the `GROUP BY` clause, the `ORDER BY` clause, and the `HAVING` clause.

Evaluation. We evaluated SQLSynthesizer's generality and accuracy in two aspects. First, we used SQLSynthesizer to solve XXX SQL exercises from a classic database textbook [22]. Textbook exercises are good resource to evaluate SQLSynthesizer's generality, since they are often designed to cover a wide range of SQL features. Some exercises are even designed on purpose to cover some less realistic, corner cases in using SQL. Second, we evaluated SQLSynthesizer on XXX SQL query related questions collected from popular online help forums, and tested whether SQLSynthesizer can synthesize correct SQL queries for them.

As a result, SQLSynthesizer successfully synthesized queries for XXX out of XXX textbook exercises and all XXX forum problems, within a very small amount of time (XXX minute per exercise or problem, on average). SQLSynthesizer's accuracy and speed make it an attractive approach to help end-users write SQL queries.

Contributions. This paper makes the following contributions:

- **Technique.** We present a technique that automatically synthesizes SQL queries from input-output examples (Section IV).
- **Implementation.** We implemented our technique in a tool, called SQLSynthesizer (Section V). It is available at: <http://sqlsynthesizer.googlecode.com>.
- **Evaluation.** We applied SQLSynthesizer to XXX textbook exercises and XXX forum questions. The experimental results show that SQLSynthesizer can synthesize a wide range of SQL queries, and it does so with small examples (Section VI).

II. ILLUSTRATING EXAMPLE

We use an example, described below, to illustrate the use of SQLSynthesizer. The example is taken from a classic database textbook [22] (Chapter 5, Exercise 1) and has been simplified for illustration purpose².

Find the name and the maximum course score of each senior student enrolled in more than 2 courses.

Despite the simplicity of the problem description, writing a correct SQL query can be non-trivial for a non-professional end-

²This exercise defines 2 tables: student and enrolled. The student table contains three columns: student_id, name, and level. Table enrolled contains three columns: student_id, course_id, and score.

student_id	name	level
1	Adam	senior
2	Bob	junior
3	Erin	senior
4	Rob	junior
5	Dan	senior
6	Peter	senior
7	Sai	senior

student_id	course_id	score
1	1	4
1	2	2
2	1	3
2	2	2
2	3	3
3	2	1
4	1	4
4	3	4
5	2	5
5	3	2
5	4	1
6	2	4
6	4	5
7	1	2
7	3	3
7	4	4

```

SELECT student.name, MAX(enrolled.score)
FROM student, enrolled
WHERE student.student_id = enrolled.student_id
      and student.level = 'senior'
GROUP BY student.student_id
HAVING COUNT(enrolled.course_id) > 2

```

name	max_score
Dan	5
Sai	5

(a) Two input tables: student (Left) and enrolled (Right)

(b) A SQL query inferred by SQLSynthesizer

(c) An output table

Fig. 1. Example input/output tables and the SQL query synthesized by SQLSynthesizer to solve the problem in Section II. In this example, users provide SQLSynthesizer with two input tables (shown in (a)) and an output table (shown in (c)). SQLSynthesizer automatically synthesizes a SQL query (shown in (b)) that transforms the two input tables into the output table.

user. Although most users can clearly understand the question, they must choose the right SQL features and use them correctly.

Users can use our SQLSynthesizer technique to obtain the desirable query. As illustrated in Figure 1, to use SQLSynthesizer, an end-user only needs to provide it with some small, representative example input and output tables (Figures 1(a) and 1(c)). Then, SQLSynthesizer works in a fully-automatic, push-button way in inferring a SQL query that satisfies the given example input and output.

The SQL query, shown in Figure 1(b), first joins two tables on the common `student_id` column, and then groups the joined result by the same `student_id` column. Further, the query selects all senior students (using a query condition in the `WHERE` clause) who has been enrolled in more than 2 courses (using a condition in the `HAVING` clause). Finally, the query projects the result on the `student.name` column and uses the `MAX` aggregate to compute the maximum course score.

III. A SQL SUBSET SUPPORTED IN SQLSYNTHESIZER

The problem of finding SQL queries satisfying a given input-output example pair is PSPACE-hard [4]. To make the problem tractable, instead of supporting all features in the standard SQL language, SQLSynthesizer focuses on a widely-used SQL subset using which a large class of query tasks can be performed. Unfortunately, when designing the SQL subset, we found that no systematic study has ever been conducted to this end, and little empirical evidence has ever been provided on which SQL features are widely-used in practice. Without such empirical knowledge, deciding which SQL subset to support remains difficult.

To address this challenge and reduce our personal bias in language design, we first conducted an online survey to ask experienced IT professionals about the most widely-used SQL features in writing database queries (Section III-A). Then, based on the survey results, we designed a SQL subset (Section III-B). We also sent the designed SQL subset to the survey participants




Fig. 2. Survey results of the most widely-used SQL features in writing a database query. There were 12 participants in the survey, and each participant was asked to select the top 10 widely-used SQL features. SQL features with no selection are omitted in this Figure for brevity.

and conducted a series of follow-up email interviews to confirm whether our design would be sufficient in practice.

A. Online Survey: Eliciting Design Requirements

Our online survey consists of 6 questions that can be divided into three parts. The first part includes simple demographic questions about participants. In the second part, participants were asked to select the top 10 most widely-used SQL features in their minds. Instead of directly asking participants about the SQL features, which might be vague and difficult to respond,

```

⟨query⟩ ::= SELECT ⟨expr⟩+ FROM ⟨table⟩+
          WHERE ⟨cond⟩+
          GROUP BY ⟨column⟩+ HAVING ⟨cond⟩+
          ORDER BY ⟨column⟩+

⟨table⟩ ::= atom
⟨column⟩ ::= ⟨table⟩.atom
⟨cond⟩ ::= ⟨cond⟩ && ⟨cond⟩
          | ⟨cond⟩ || ⟨cond⟩
          | ( ⟨cond⟩ )
          | ⟨cexpr⟩ ⟨op⟩ ⟨cexpr⟩
⟨op⟩ ::= = | > | <
⟨cexpr⟩ ::= const | ⟨column⟩
⟨expr⟩ ::= ⟨cexpr⟩ | COUNT(⟨column⟩) | COUNT(DISTINCT ⟨column⟩)
          | SUM(⟨column⟩) | MAX(⟨column⟩) | MIN(⟨column⟩)

```

Fig. 3. Syntax of the supported SQL subset in SQLSynthesizer: *const* is a constant value and *atom* is a string value, representing a table name or a column name.

we presented them a list of *all* standard SQL features in writing a query. Additionally, participants were asked to report their own experience in writing SQL queries in the third part of the survey.

We sent out invitation to the graduate mailing list at University of Washington, and posted our survey on professional online forums (e.g., StackOverflow). As of April 2013, we received 12 responses. On average, the respondents have 9.5 years of experience in software development (max: 15, min: 5), and 5.5 years of experience in using database (max: 10, min: 2). In addition, two participants identified themselves as database professionals.

Figure 2 summarizes the survey results.

B. Language Syntax

Based on the survey results, we design a SQL subset whose syntax is shown in Figure 3. Its semantics are the same as the standard SQL language, and are omitted in this section for brevity.

The supported SQL subset is a subset of the standard SQL language. It covers all top 10 most widely-used SQL features voted by the survey participants, except for the `IN` keyword in Figure 2. In addition, the SQL subset supports the `HAVING` keyword since `HAVING` is often used together with the `GROUP BY` clause. The SQL subset, though by means complete in writing all possible queries, has significantly enriched the SQL subset supported by the existing query inference work [4], [26]. Besides supporting the standard select-from-where queries as in [4], [26], our SQL subset also supports table joins, aggregates (i.e., `COUNT`, `MAX`, `MIN`, and `AVG`), the `GROUP BY` clause, the `ORDER BY` clause, and the `HAVING` clause. For readers who are not familiar with the basic SQL idioms, we show an example query using our SQL subset in Figure 4, and annotate it with some basic concept.

When designing this SQL subset, we only considered standard SQL features, while excluding user-defined functions and vendor-specific features, e.g., the `TOP` keyword supported

```

          Projection column   Aggregate
                ↓              ↓
SELECT student.name, MAX(enrolled.score)
FROM student, enrolled ← Query tables
WHERE student.student_id = enrolled.student_id ← Join conditions
      and student.level = 'senior' ← Query conditions
GROUP BY student.student_id ← GROUP BY clause
HAVING COUNT(enrolled.course_id) > 2 ← HAVING clause
ORDER BY student.name ← ORDER BY clause

```

Fig. 4. An example query using the SQL subset defined in Figure 3.

in Microsoft SQLServer. We discarded some standard SQL features, primarily for three reasons. First, some features are designed as syntactic sugar to make a SQL query easier to write; and thus can be safely removed without affecting a language’s functionality. For example, the `BETWEEN` keyword checks whether a given value is within a specific range, and can be simply replaced by two query conditions. Similarly, the `NOT NULL` keyword is also omitted. Second, some features, such as `FULL JOIN`, `LEFT JOIN`, and `RIGHT JOIN`, provide special ways to join tables, and are less likely to be used by non-expert end-users. Third, other features, such as `IN`, `UNION`, and `NOT EXIST`, are used to write sub-queries, which are the major source of the PSPACE-hardness in inferring a SQL query [4]. We exclude them in order to make the synthesis problem more tractable. The `LIKE` is also discarded, since it is used for string wildcard matching, and can also lead to **[[xx]]**.

C. Follow-up Interviews: Feedback about the SQL Subset

After proposing the SQL subset in Figure 3, we performed follow-up email interviews to gain participants’ feedback about it. Participants were first asked to rate the expressiveness of the SQL subset in Figure 3 in writing real-world database queries, on a 6-point scale (5-completely sufficient; 0-not sufficient at all; and in-between values indicating intermediate sufficiency), and then to provide their comments.

On average, the rating of this SQL subset is 4.5. Most of the participant rated it 5, or 4. Only one participant rated it 3, because this participant misinterpreted the designed syntax and thought it does not support table joins.

Overall, based on the feedback by experienced IT professionals, we think our SQL subset is reasonably expressive for writing database queries that most end-users need.

IV. TECHNIQUE

This section first gives an overview of SQLSynthesizer’s workflow and high-level algorithm in Section IV-A, and then explains SQLSynthesizer’s three steps in details (Section IV-B, Section IV-C1, and Section IV-D).

A. Overview

Figure 5 illustrates SQLSynthesizer’s workflow. SQLSynthesizer consists of three steps: (1) the “Skeleton Creation” step (Section IV-B) infers a set of query skeletons from the given examples; (2) the “Query Completion” step (Section IV-C) infers the missing parts in each query skeleton and outputs




Fig. 5. Illustration of SQLSynthesizer’s workflow of synthesizing SQL queries from input-output examples.

Input: example input table(s) T_I , example output table T_O

Output: a ranked list of SQL queries
 $\text{synthesizeSQLQueries}(V_{old}, V_{new}, T)$

```

1:  $queryList \leftarrow$  an empty list
2:  $skeletons \leftarrow \text{createQuerySkeletons}(T_I, T_O)$ 
3: for each  $skeleton$  in  $skeletons$  do
4:    $conds \leftarrow \text{inferConditions}(T_I, T_O, skeleton)$ 
5:    $aggs \leftarrow \text{searchForAggregates}(T_I, T_O, skeleton, conds)$ 
6:    $columns \leftarrow \text{searchForOrderBys}(T_O, skeleton, aggs)$ 
7:    $queries \leftarrow \text{buildQueries}(skeleton, conds, aggs, columns)$ 
8:   for each  $query$  in  $queries$  do
9:     if  $\text{isValidOnExamples}(query, T_I, T_O)$  then
10:       $queryList.add(query)$ 
11:     end if
12:   end for
13: end for
14:  $\text{rankQueries}(queryList)$ 
15: return  $queryList$ 

```

Fig. 6. Algorithm for synthesizing SQL queries from input-output examples.

a list of syntactically-valid queries that satisfy the provided example input and output; and (3) the “Candidate Ranking” step (Section IV-D) ranks all synthesized SQL queries and place the more likely ones near the top. Users can inspect the query list and select a query from it. If SQLSynthesizer produces SQL queries that satisfy the input and output examples, but does not address the intention that the user wants; SQLSynthesizer can be used interactively by asking users to give more informative examples and then refine the inferred queries.

Figure 6 sketches SQLSynthesizer’s high-level algorithm. Line 2 corresponds to the first “Query Skeleton Creation” step. Lines 3 – 13 correspond to the second “Query Completion” step, in which SQLSynthesizer searches for the query conditions (line 4), aggregates (line 5), and columns in the `ORDER BY` clause (line 6). SQLSynthesizer then assembles a list of candidate SQL queries (line 7), and validates their correctness on the examples (lines 8 – 11). Line 14 corresponds to the “Query Ranking” step.

B. Skeleton Creation

A query skeleton is a partially-complete SQL query. It consists of three parts: the tables used in the query, the table columns used to join tables, and the table columns used to project the query results. Other elements in the query, such as query conditions and aggregate functions, will be determined in the next step.

A query skeleton captures the most basic structure of a

SQL query; creating it is the first step of synthesizing a complete query. SQLSynthesizer performs a simple scan over the examples, and employs several heuristics to determine the table set, joining columns, and project columns.

[[why heuristics. the source of PSPACE]]

Determining the Table Set. We observe that End-users are often unwilling to provide more than enough input: every table in the example input is expected to be used (at least once) in the result query. Thus, we assume that every input table should be used in the query. On the other hand, it is possible that one input table will be used for multiple times in a query. SQLSynthesizer does not forbid this case, rather, it uses a heuristic to estimate the table set: if one column from an input table appears multiple times in the example output, we add the input table to the table set the same number of times. [[xx]]

The rationale behind this heuristic is that if a table column appears multiple times in the output, it may indicate that the table would be joined multiple times.

Determining the Joining Columns. Given a set of tables, there are many ways to join them. Enumerating all possibilities leads to a huge number of joining conditions and would quickly become intractable. To make it feasible, we use three simple but effective rules to capture the most likely ways to join tables in practice. First, tables are often joined on their primary keys with the same data type. For example, in Figure 1, the `student` table can be joined with the `enrolled` table on the `student_id` column. By contrast, it is unlikely to join two tables on columns with different types. Second, tables are often joined on columns with the same name, since columns with the same name often such as joining the `student` table with the `enrolled` table on the `student_name` column. Third, it is only meaningful to join two tables on columns that have the same data type and some overlapped values. [[revise]]

SQLSynthesizer restricts the search space in uses the above three rules [[need to revise]]

[[need to implement above.]] [[mention how many skeletons will be created]] [[give an algorithm]]

Determining the Output Columns. For each column in the output table, SQLSynthesizer first checks whether its column name appears in any of the input table. If so, SQLSynthesizer uses the matched column from the input table as the output column. Otherwise, the output column must be produced by using an aggregate function. [[same names]] Consider the example in Figure 1, SQLSynthesizer determines that column `name` comes from the `student` table, while column `max_Score` must be created by using an aggregation operator. [[If there


```

SELECT      student.name, <Aggregate>
FROM        student, enrolled
WHERE       student.student_id = enrolled.student_id
           and <Query Condition>
GROUP BY    < Column Name(s)>
HAVING      <Query Condition>
ORDER BY    <Column Name(s)>

```

Fig. 7. A query skeleton created for the motivating example in Figure 1.

is no column name]] [[check the values in the output column]]

[[It is possible that multiple skeleton can be created. add an algorithm here.]]

Figure 7 shows the created query skeleton for the motivating example in Figure 1. In this skeleton, three unknown structures represented by <Aggregation> or <Conditions> are in red, and will be filled in the next phase. [[revise text]]

[[how to create group by]]

C. Query Completion

In this step, SQLSynthesizer takes as input a query skeleton, and infers the query conditions (Section IV-C1), aggregates (Section IV-C2), and order-by clauses (Section IV-C3); and finally produces a set of syntactically-correct SQL queries.

1) *Inferring Query Conditions*: SQLSynthesizer casts the problem of *query condition inference* as finding appropriate *rules* that can perfectly divide the whole searching space into positive part and negative part. In our context, the search space contains all tuples generated by joining the input tables; the positive part are all tuples in the output table; and the negative part are the rest tuples.

A natural way to learn rules is using a decision-tree-based learning algorithm to infer a set of rules as query conditions. However, a key challenge to effectively employ the decision tree algorithm is how to devise expressive features. Existing approaches [26] simply use data values in each tuple as features. Although using data values as features is sufficient for simple query conditions like “student.level = ‘senior’”, it suffers from losing much useful *structure information* needed in a SQL query, and fails to infer complex query conditions like “count(enrolled.course_id) > 3 (after grouping by the student.student_id column)”.

SQLSynthesizer addresses this challenge by enriching the existing data value features with two types of additional features:

- **Aggregation Features**. SQLSynthesizer groups the data values in a table by each column, and applies every aggregate functions (i.e., COUNT, MAX, MIN and AVG) to compute the values for the other columns. [[cannot infer group by column1, column2]] [[add an example]]
- **Comparison Features**. For each tuple, SQLSynthesizer compares the data values of every two type-comparable columns, and records the comparison results (1 or 0) as features. [[add an example]]

The above two additional features seamlessly encodes SQL structure knowledge encoding permits our technique to make

use of correlations between columns, rather than only values from each isolated and sequential columns. Table 8 shows an example.

[[The increasing number of features, can be falsified quickly]]

After enriching the original data values with additional features, SQLSynthesizer uses a variant of the decision tree algorithm, called PART [8], to infer a set of accurate and expressive rules. Compared to the original decision tree algorithm [], PART has two notable features. [[unclear]] First, it uses a “divide-and-conquer” strategy to repeatedly construct rules and remove the tuples that have been covered until no tuples are left, and thus is more efficient. Second, when constructing each rule, PART uses a pruned decision tree built from current set of tuples and only makes the leaf with the largest coverage into the resulting rules, without keeping the whole learned tree in memory. This permits PART [[what good results?]] and is more suitable for the SQL inference problem.

Take the joined table in Figure 1 as an example, SQLSynthesizer infers the following rules to perfectly divide the joined table into the query output:

```

COUNT(enrolled.course_id) > 2 (group by student_id )
&& student.level = 'senior'.

```

SQLSynthesizer next splits the inferred rules and put each condition into the appropriate places in a SQL query (based on the SQL language’s syntax). For the example in Figure 1, SQLSynthesizer extracts && student.level = ‘senior’ as the query condition, treats student_id as the group by column, and uses COUNT(enrolled.course_id) > 2 in the Having clause.

2) *Searching for Aggregates*: For each column produced by an aggregation operator, the whole search space includes all possible combinations of table columns and the five supported aggregation operators (see Figure 3). SQLSynthesizer leverages the following two observations to further reduce the search space:

- The data type of an output column must be compatible with the aggregation operator’s return type. For instance, if an output column has the String type, it must not use aggregation operators (e.g., count and sum) that returns an Integer.
- If an arithmetic aggregation operator, such as max and min, is used, each value in the output column must has appeared in the input table.

3) *Searching for Order By Columns*: After the query conditions and aggregates are inferred, SQLSynthesizer observes data values in each output table column. If the data values in a column appear are sorted, SQLSynthesizer append the column name to the Order By clause.

D. Candidate Ranking

It is possible that multiple SQL queries satisfying the given input-output examples will be returned. This may adversely impact end-users who want to perform simple query tasks but now need to select the query of their intent. To alleviate this problem, we employ the Occam’s razor principle, which states

An input table		Aggregation Features										Comparison Features		
C1	C2	Group by C1					Group by C2					C1 = C2	C1 < C2	C1 > C2
2	4	Count	Count Distinct	Min	Max	Avg	Count	Count Distinct	Min	Max	Avg	0	1	0
2	1	3	2	1	4	2	1	1	2	2	2	0	0	1
2	1	3	2	1	4	2	3	2	1	2	5/3	0	0	1
1	1	3	2	1	4	2	3	2	1	2	5/3	0	0	1
		4	2	1	1	1	3	2	1	2	5/3	1	0	0

Fig. 8. Illustration of the aggregation features and the comparison features enriched by SQLSynthesizer. (Left) An example input table with two columns: C1 and C2. (Center) The aggregation features enriched by SQLSynthesizer for the input table. (Right) The comparison features enriched by SQLSynthesizer for the input table.

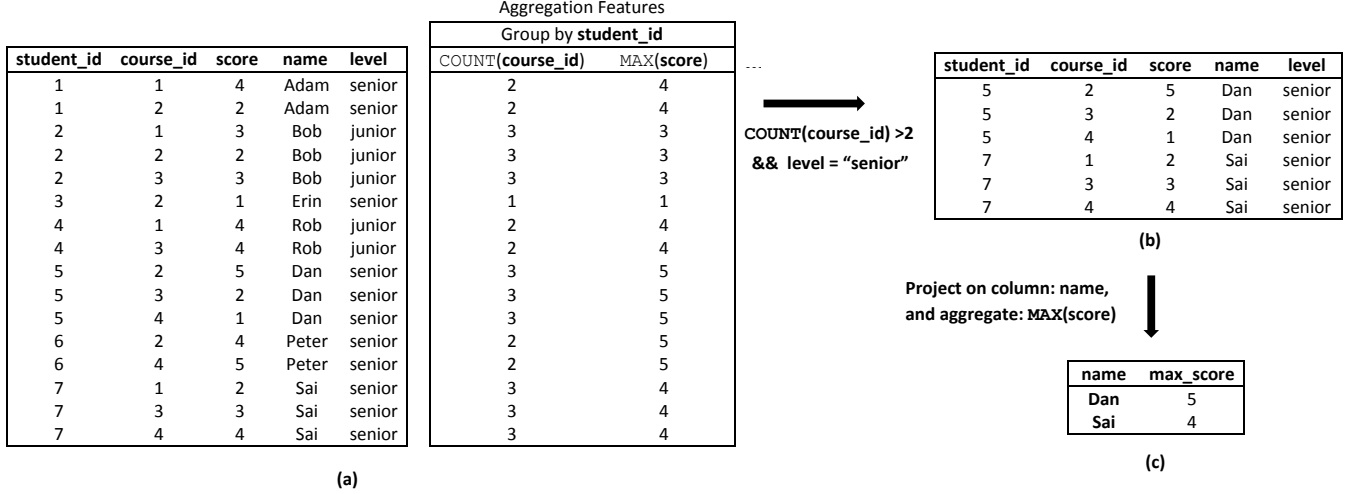


Fig. 9. Illustration of how additional features added by SQLSynthesizer helps in inferring query conditions. (a) shows SQLSynthesizer enriches the original table (Left: the result of joining table student with table enrolled on the student_id column) with additional features. For brevity, only relevant aggregation features are shown. Using the added aggregation features, SQLSynthesizer infers two query conditions that transform the original table into the table show in (b). (c) shows the output table, which is produced by projecting the table in (b) on its column: name, and an aggregate: MAX(score).

that the simplest explanation is usually the correct one, to rank a more likely query higher in the output list. A simpler query is less likely to overfit the given examples than a complex query, even when both of them can transform the example input to the example output.

A SQL query is simpler than another one if it uses fewer query conditions (including conditions in the `Having` and `from` clauses) or the expressions (including aggregates) in each query condition are pairwise simpler (e.g., expression `Count(student_id)` is simpler than `Count(Distinct student_id)`). Simpler query conditions suggests the extraction logics are more common and general.

In our implementation, SQLSynthesizer computes a cost for each query, and prefers queries with lower costs. The cost for a SQL query is computed approximately by summarizing the number of conditions, aggregates, and other expressions appearing in the `group by` and `order by` clauses. This heuristic, though fairly simple, has been observed to work well. Figure 10 shows an example.

E. Discussion

Soundness and Completeness. The SQLSynthesizer technique is neither sound nor complete. The primary reason is that several steps (e.g., the Query Skeleton Creation step in Section IV-B) use heuristics to infer possible joins, tables used in the `from` clause, and columns in the `select` clause.

name	score
Bob	4
Dan	5
Jim	2

(a) The input table: student

(b) The output table

1. `select name from student where score > 2`
2. `select name from student where name = 'Bob' or name = 'Dan'`

Fig. 10. Illustration of SQLSynthesizer's query candidate ranking heuristic. SQLSynthesizer produces two queries for the given input-output examples. Based on the heuristic in Section IV-D, the first query differs from the second query by using simpler conditions, and thus is ranked higher.

Such heuristics are necessary, since they provide a good approximate solution to the problem of finding SQL queries from examples, which has been proved to be PSPACE-hard and is thus intractable in practice. Although SQLSynthesizer cannot guarantee to infer correct SQL queries for all cases, as demonstrated in Section VI, we find SQLSynthesizer is useful in synthesizing a wide variety of queries in practice.

[[no title]]

V. IMPLEMENTATION

We implemented the proposed technique in a tool, called SQLSynthesizer. SQLSynthesizer uses the built-in PART al-

gorithm implementation in the Weka toolkit [13] to learn query conditions (Section IV-C). SQLSynthesizer also uses MySQL [21] as the backend database to validate the correctness of each synthesized SQL query. Specifically, SQLSynthesizer first populates the backend database with the given input tables; when a SQL query is synthesized, SQLSynthesizer executes the query on the database to observe whether the output matches the given output. Our implementation is publicly available at: <http://sqlsynthesizer.googlecode.com>

VI. EVALUATION

We evaluated four aspects of SQLSynthesizer’s effectiveness, answering the following research questions:

- What is the success ratio of SQLSynthesizer in synthesizing SQL queries? (Section VI-C1).
- How long does it take for SQLSynthesizer to synthesize a SQL query (Section VI-C2).
- How much human effort is needed to write sufficient input-output examples for SQL synthesis (Section VI-C3).
- How does SQLSynthesizer’s effectiveness compare to existing SQL query inference techniques (Section VI-C4).

A. Benchmarks

We collected benchmarks from two sources:

- We selected *all* SQL query related exercises (XXX in total) from a classic database textbook [22]. All exercises are from Chapter 5, which systematically introduces the SQL language. Textbook exercises are good resource to evaluate SQLSynthesizer’s generality, since such exercises are often designed to cover a wide range of SQL features. Some exercises are even designed on purpose to cover some less realistic, corner cases in using SQL. As shown in Figure 11, each textbook exercises involves at least 3 tables. It was unintuitive for us to write the correct query by simply looking at the problem description in the exercise.
- We searched SQL query related questions raised by real-world database users from 3 popular online forums [5], [25], [27]. We focused on questions about using standard SQL features rather than vendor-specific SQL features. We excluded questions that were vaguely described or obviously wrong, and discarded questions that had been proved to be unsolvable by using SQL (e.g., computing a transitive closure). We collected XXX recent forum questions related to writing a SQL query. **[[merge same types. exclude jdbc, why relative few]]**

B. Evaluation Procedure

We used SQLSynthesizer to solve each textbook exercise and forum question. If an exercise or problem was associated with example input and output, we directly applied SQLSynthesizer on those examples. Otherwise, we manually wrote some example input and output. To reduce the bias in writing examples, all examples are written by a different graduate student (whose research field is not database-related)

from University of Washington other than SQLSynthesizer’s developers.

We checked SQLSynthesizer’s correctness by comparing its output with the expected SQL queries. Specifically, for textbook exercises, we compared SQLSynthesizer’s output with their correct answers; for forum questions, we manually wrote the correct SQL query and then determined whether SQLSynthesizer can produce it.

For some textbook exercises and forum questions, SQLSynthesizer inferred a SQL query that satisfied the input-output examples, but did not behave as we expected when applied to other inputs. We manually found another input on which the SQL query mis-behaved and re-applied SQLSynthesizer to the new input. We repeated this process and recorded the number of interactions until SQLSynthesizer synthesized a desirable SQL query.

All experiments were run on a 2.67GHz Intel Core PC with 4GB physical memory (2GB was allocated for the JVM), running Windows 7.

C. Results

Figure 11 summarizes our experimental results.

1) *Success Ratio*: As shown in Figure 11, SQLSynthesizer synthesized expected SQL queries for XXX out of XXX the textbook exercises, and XXX out of XXX the forum questions.

[[why the technique can work, why some problem can not be solved]]

2) *Performance*: We measured SQLSynthesizer’s performance by recording the average time cost in producing a ranked list of SQL queries. As shown in Figure 11, the performance of SQLSynthesizer is reasonable. On average, it uses less than XXX minutes to produce the results in one interactive round. Most of the time is spent querying the backend database to validate the correctness of each synthesized SQL query.

3) *Human Efforts*: We measured the human efforts taken to use SQLSynthesizer in two ways. First, the time cost to write input-output examples. Second, the number of interactive rounds in invoking SQLSynthesizer to synthesize the desirable SQL queries.

As shown in Figure 11, human efforts spent in providing input-output examples are very limited: on average, it took less than 5 minutes for one benchmark. **[[explain some abnormal points]]**

The number of interactive rounds is a measure of the generalization power of the conditional learning part of the algorithm and the ranking scheme. We observed that the tool typically requires just XXX rounds of interaction, when the user is smart enough to give an example for each input format (which typically range from 1 to 3) to start with. This was indeed the case for most cases in our benchmarks, even though our algorithm can function robustly without this assumption. The maximum number of interactive rounds required in any scenario was **[[XXX]]** (with 2 to 3 being a more typical number). **[[the largest table]]** The maximum number of examples required in any scenario over all possible interactions was 10.

forum questions, though covering a wide variety of SQL features, may not be representative enough. Thus, we can not claim the results can be generalized to an arbitrary use-case scenario. Second, we only compared SQLSynthesizer with the *Query by Output* technique [26]. Using other query inference or recommendation techniques might achieve different results. Third, our experiments focus on evaluating SQLSynthesizer’s generality and accuracy. Even though all experiments are carried out by a different person other than SQLSynthesizer’s developers, it is unknown about SQLSynthesizer’s general usability in practice. To address this issue, we plan to conduct a user study in our future work.

Experimental Conclusions. We have three chief findings: (1) The supported SQL subset in SQLSynthesizer is expressive enough to describe a variety of database queries. (2) SQLSynthesizer can efficiently synthesize desirable SQL queries with a small amount of human efforts and small input-output examples. (3) SQLSynthesizer produces better results than an existing technique (*Query by Output* [26]).

VII. RELATED WORK

This section discusses two categories of closely-related work on reverse engineering SQL queries and automated program synthesis.

A. Reverse Engineering SQL Queries

Reverse engineering SQL queries is a well-known technique in the database community [4], [26], [28] to enhance a database system’s usability. Zloof’s pioneering work on *Query by Example* (QBE) [28] provided a high-level query language and a form-based Graphical User Interface (GUI) for writing database queries. To use the QBE system, users need to learn its own query language, formulate a query with the language, and fill in the appropriate skeleton tables on the GUI. By contrast, SQLSynthesizer mitigates such learning curves by only requiring users to provide some representative examples to describe their query intentions

Tran et al. [26] proposed a technique, called *Query by Output* (QBO), to find a set of semantically-equivalent SQL queries to a given input query. QBO is related to but significantly differs from SQLSynthesizer in two aspects. First, QBO has a rather different goal and requires different inputs : it takes as inputs a database, an SQL query, and the query’s output (on the database); and computes one or more equivalent queries that produce the same output on the input database. Second, QBO can only infer simple select-project-join queries, while excluding many useful SQL features, such as aggregates, the *Having* clause, and the *Group by* clause. As we demonstrated in in Section VI, queries inferred by QBO cannot be applied to many of the other real-world cases. QBO’s key limitation stems from the fact that it only considers existing tuple values in the input tables as features, when learning a set of classification rules as query conditions. By contrast, SQLSynthesizer remedies this limitation by enhancing the existing tuple values with two kinds of additional features (Section IV-C1).

Recently, Sarma et al. [4] studied the *View Definitions Problem* (VDP). VDP aims to find the most succinct and accurate view definition, when the view query is restricted to a specific family of queries. VDP can be solved as a special case in SQLSynthesizer where there is only one input table and one output table. Furthermore, the main contribution of Sarma et al’s work is the complexity analysis of three variants of the view definitions problem; there is no tool implementation or empirical studies to evaluate the proposed technique.

B. Automated Program Synthesis

Program synthesis [10] is a useful technique to create an executable program in some underlying language from specifications that can range from logical declarative specifications to examples or demonstrations [1], [2], [6], [11], [14], [16], [18], [19], [23]. It has been used recently for many applications such as synthesis of efficient low-level code [24], data structure manipulations [7], geometry constructions [12], snippets of excel macros [14], relational data representations [1], [2] and string expressions [11], [23].

The PADS [7] system takes a large sample of unstructured data and infers a format that describes the data. Related, the Wrangler tool, developed in the HCI community, provides a visual programming-by-demonstration interface to table transformations for data cleaning [16]. These two techniques, though well-suited for tasks like text extraction, are inapplicable to synthesizing a database query, since they use completely different abstractions than SQL and lack the support for many database operations like table joining, aggregations, etc.

Harris and Gulwani described a system for learning excel spreadsheet transformation macros from an example input-output pair [14]. Given one input table and one output table, their system can infer an excel macro that filters, duplicates, and/or reorganizes table cells to generate the output table. SQLSynthesizer differs in multiple respects. First, excel macros have significantly different semantics than the SQL language. An excel macro can express a variety of table transformation operations (e.g., table re-shaping), but are not capable to formulate database queries. Second, Harris and Gulwani’s approach treats table cells as atomic units, and thus has different expressiveness than SQLSynthesizer. For instance, their technique can generate macros to transform one table to another, but cannot join multiple tables or group query results by certain table columns.

Some recent work proposed query recommendations systems to reduce the obstacles to using relational databases [15], [17]. SQLShare [15] is a cloud-based service that allows users to upload their data and SQL queries. Each uploaded query is saved as a view, allowing other users to compose and reuse. SnipSuggest [17] is a SQL autocompletion system. As a user types a query, SnipSuggest mines existing query logs to recommend relevant clauses or SQL snippets (e.g., the table names for the *from* clause) based on the partial query that the user has typed so far. Compared to SQLSynthesizer, both SQLShare and SnipSuggest assume the existence of a comprehensive query log (produced either by the user herself

or other users) that contains valuable information. However, such assumption often does not hold for many database users in practice. SQLSynthesizer eliminates this assumption and infers SQL queries by using user-provided examples.

Cheung et al [3] presented a technique to infer SQL queries from imperative code. Their technique identifies fragments of application logic (written in an imperative language like Java) that can be pushed into SQL queries. Compared to SQLSynthesizer, their work is designed for developers to improve a database application's performance, rather than helping non-expert end-users write correct SQL queries from scratch. If an end-user wishes to use their technique to synthesize a SQL query, she must write a snippet of imperative code to describe the query task. Comparing to providing example input and output as required by SQLSynthesizer, writing correct imperative code can be too challenging for a typical end-user, and thus would substantially degrade the technique's usability in practice.

VIII. CONCLUSION AND FUTURE WORK

This paper studied the problem of automated SQL query synthesis from simple input-output examples, and presented a practical technique (and its tool implementation, called SQLSynthesizer). SQLSynthesizer is motivated by the growing population of non-expert database users, who need to query on their databases, but have difficulty with SQL. We have shown that SQLSynthesizer is able to synthesize a variety of SQL queries, and it does so with small input-output examples. We view SQLSynthesizer as an important step toward making databases more usable. The source code of our tool implementation is available at: <http://sqlsynthesizer.googlecode.com>

For future work, we are interested in conducting a user study to evaluate SQLSynthesizer's usability. We also plan to explore applications of this technique.

REFERENCES

- [1] A. Arasu, S. Chaudhuri, and R. Kaushik. Learning string transformations from examples. *Proc. VLDB Endow.*, 2(1):514–525, Aug. 2009.
- [2] D. M. Barbosa, J. Cretin, N. Foster, M. Greenberg, and B. C. Pierce. Matching lenses: alignment and view update. In *Proceedings of the 15th ACM SIGPLAN international conference on Functional programming*, ICFP '10, pages 193–204, New York, NY, USA, 2010. ACM.
- [3] A. Cheung, A. Solar-Lezama, and S. Madden. Inferring sql queries using program synthesis. *CoRR*, abs/1208.2013, 2012.
- [4] A. Das Sarma, A. Parameswaran, H. Garcia-Molina, and J. Widom. Synthesizing view definitions from data. In *Proceedings of the 13th International Conference on Database Theory*, ICDT '10, pages 89–103, New York, NY, USA, 2010. ACM.
- [5] Database Journal. <http://forums.databasejournal.com/>.
- [6] K. Fisher. Learnpads: Automatic tool generation from ad hoc data. In *In SIGMOD*, 2008.
- [7] K. Fisher, D. Walker, K. Q. Zhu, and P. White. From dirt to shovels: fully automatic tool generation from ad hoc data. In *Proceedings of the 35th annual ACM SIGPLAN-SIGACT symposium on Principles of programming languages*, POPL '08, pages 421–434, New York, NY, USA, 2008. ACM.
- [8] E. Frank and I. H. Witten. Generating accurate rule sets without global optimization. In *Proceedings of the 15th International Conference on Machine Learning*, ICML'98, pages 144–151. Morgan Kaufmann, 1998.
- [9] J. Gray, D. T. Liu, M. Nieto-Santisteban, A. Szalay, D. J. DeWitt, and G. Heber. Scientific data management in the coming decade. *SIGMOD Rec.*, 34(4):34–41, Dec. 2005.
- [10] S. Gulwani. Dimensions in program synthesis. In *Proceedings of the 12th international ACM SIGPLAN symposium on Principles and practice of declarative programming*, PPDP '10, pages 13–24, New York, NY, USA, 2010. ACM.
- [11] S. Gulwani. Automating string processing in spreadsheets using input-output examples. In *Proceedings of the 38th annual ACM SIGPLAN-SIGACT symposium on Principles of programming languages*, POPL '11, pages 317–330, New York, NY, USA, 2011. ACM.
- [12] S. Gulwani, V. A. Korthikanti, and A. Tiwari. Synthesizing geometry constructions. In *Proceedings of the 32nd ACM SIGPLAN conference on Programming language design and implementation*, PLDI '11, pages 50–61, New York, NY, USA, 2011. ACM.
- [13] M. Hall, E. Frank, G. Holmes, B. Pfahringer, P. Reutemann, and I. H. Witten. The weka data mining software: an update. *SIGKDD Explor. Newsl.*, 11(1):10–18, Nov. 2009.
- [14] W. R. Harris and S. Gulwani. Spreadsheet table transformations from examples. In *Proceedings of the 32nd ACM SIGPLAN conference on Programming language design and implementation*, PLDI '11, pages 317–328, New York, NY, USA, 2011. ACM.
- [15] B. Howe, G. Cole, N. Khoussainova, and L. Battle. Automatic example queries for ad hoc databases. In *Proceedings of the 2011 international conference on Management of data*, SIGMOD '11, pages 1319–1322, New York, NY, USA, 2011. ACM.
- [16] S. Kandel, A. Paepcke, J. Hellerstein, and J. Heer. Wrangler: interactive visual specification of data transformation scripts. In *Proceedings of the 2011 annual conference on Human factors in computing systems*, CHI '11, pages 3363–3372, New York, NY, USA, 2011. ACM.
- [17] N. Khoussainova, Y. Kwon, M. Balazinska, and D. Suciu. Snipsuggest: context-aware autocompletion for sql. *Proc. VLDB Endow.*, 4(1):22–33, Oct. 2010.
- [18] T. Lau, S. A. Wolfman, P. Domingos, and D. S. Weld. Programming by demonstration using version space algebra. *Mach. Learn.*, 53(1-2):111–156, Oct. 2003.
- [19] T. A. Lau, P. Domingos, and D. S. Weld. Version space algebra and its application to programming by demonstration. In *Proceedings of the Seventeenth International Conference on Machine Learning*, ICML '00, pages 527–534, San Francisco, CA, USA, 2000. Morgan Kaufmann Publishers Inc.
- [20] MDX: A query language for OLAP databases. <http://msdn.microsoft.com/en-us/library/gg492188.aspx>.
- [21] MySQL. <http://www.mysql.com>.
- [22] R. Ramakrishnan and J. Gehrke. *Database Management Systems*. Addison-Wesley (3rd Edition), 2007.
- [23] R. Singh and S. Gulwani. Learning semantic string transformations from examples. In *Proceedings of the 37th International Conference on Very Large Data Bases*, VLDB '12, New York, NY, USA, 2012. ACM.
- [24] A. Solar-Lezama, R. Rabbah, R. Bodík, and K. Ebcioglu. Programming by sketching for bit-streaming programs. In *PLDI*, 2005.
- [25] StackOverflow. <http://www.stackoverflow.com>.
- [26] Q. T. Tran, C.-Y. Chan, and S. Parthasarathy. Query by output. In *Proceedings of the 35th SIGMOD international conference on Management of data*, SIGMOD '09, pages 535–548, New York, NY, USA, 2009. ACM.
- [27] Tutorialized Forums. <http://forums.tutorialized.com>.
- [28] M. M. Zloof. Query-by-example: the invocation and definition of tables and forms. In *Proceedings of the 1st International Conference on Very Large Data Bases*, VLDB '75, pages 1–24, New York, NY, USA, 1975. ACM.